

# Coursera-Stanford-ML-Notes

Quentin Truong  
Taught by Professor Andrew Ng

20 June 2017 - 21 July 2017

## Contents

<b>1</b>	<b>Week 1: Introduction</b>	<b>2</b>
1.1	Overview . . . . .	2
<b>2</b>	<b>Week 2: Linear Regression with Multiple Variables</b>	<b>3</b>
2.1	Overview . . . . .	3
2.2	Notation . . . . .	3
2.3	Gradient Descent . . . . .	3
2.4	Normal Equation . . . . .	3
<b>3</b>	<b>Week 3: Logistic Regression</b>	<b>4</b>
3.1	Overview . . . . .	4
3.2	Logistic Regression Hypothesis Function . . . . .	4
3.3	Logistic Regression Cost Function . . . . .	4
3.4	Proof of Logistic Regression Cost Function Derivative . . . . .	5
3.5	Regularization . . . . .	6
<b>4</b>	<b>Week 4: Artificial Neural Networks Representation</b>	<b>7</b>
4.1	Overview . . . . .	7
4.2	ANN Notation . . . . .	7
4.3	ANN Equations . . . . .	7
4.4	Sample Three Layer System . . . . .	8
<b>5</b>	<b>Week 5: Artificial Neural Network Learning</b>	<b>9</b>
5.1	Notation . . . . .	9
5.2	Cost Function . . . . .	9
5.3	Backpropagation Algorithm . . . . .	9
5.4	Backpropagation Derivation - Base Case . . . . .	10
5.5	Backpropagation Derivation - Recursive Case . . . . .	10
5.6	Backpropagation Intuition . . . . .	11
5.7	Gradient Checking . . . . .	11
5.8	Random Initialization . . . . .	11
<b>6</b>	<b>Week 6: Advice for Machine Learning</b>	<b>12</b>
6.1	Overview . . . . .	12
6.2	Evaluating Hypothesis . . . . .	12
6.3	Bias, Variance, Learning Curves . . . . .	12
6.4	Error Analysis . . . . .	13
6.5	Precision and Recall . . . . .	13
6.6	Large Data Sets . . . . .	13

<b>7</b>	<b>Week 7: Support Vector Machines</b>	<b>14</b>
7.1	Overview . . . . .	14
7.2	Large Margin Classification . . . . .	14
7.3	Math behind SVM . . . . .	14
7.4	Kernels . . . . .	14
<b>8</b>	<b>Week 8: Unsupervised Learning</b>	<b>15</b>
8.1	Overview . . . . .	15
8.2	K-Means Notation . . . . .	15
8.3	K-Means Distortion Function/Optimization Objective . . . . .	15
8.4	K-Means Algorithm . . . . .	15
8.5	Dimensionality Reduction: Principal Component Analysis . . . . .	15
8.6	PCA Notation . . . . .	15
8.7	PCA Equations . . . . .	15
8.8	PCA Algorithm . . . . .	16
8.9	PCA K . . . . .	16
<b>9</b>	<b>Week 9: Anomaly Detection</b>	<b>17</b>
9.1	Overview . . . . .	17
9.2	Gaussian (Normal) Distribution Equations . . . . .	17
9.3	Anomaly Detection Algorithm . . . . .	17
9.4	Multivariate Gaussian Distribution Equations . . . . .	17
9.5	Recommender Systems Notation . . . . .	17
9.6	Recommender Systems Algorithm . . . . .	18
9.7	Collaborative Filtering Equations . . . . .	18
9.8	Collaborative Filtering Algorithm . . . . .	18
<b>10</b>	<b>Week 10: Large Scale ML</b>	<b>19</b>
10.1	Overview . . . . .	19
<b>11</b>	<b>Week 11: Photo OCR</b>	<b>20</b>
11.1	Overview . . . . .	20
11.2	Summary . . . . .	20

# 1 Week 1: Introduction

## 1.1 Overview

- Machine Learning: “A computer program is said to learn from experience  $E$  with respect to some class of tasks  $T$  and performance measure  $P$ , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ .”
- Supervised Learning: know what our correct output looks like
  - Regression: want continuous output
  - Classification: want discrete output
- Unsupervised Learning: little or no idea what our results should look like
  - Clustering: find groups according to similarity in various variables
  - Nonclustering: find structure in chaos

## 2 Week 2: Linear Regression with Multiple Variables

### 2.1 Overview

- Use linear regression for continuous output
- Choose gradient descent if many features (million+) because the inverse matrix required for the normal equation can become expensive to compute
- Normal equation will directly compute theta
- Normalize features if using gradient descent

### 2.2 Notation

$m = \text{number of samples}$

$n = \text{number of feature}$

$x = (n \times 1)$

$X = (m \times n)$

$X_j = (m \times 1)$

$\theta = (n \times 1)$

$\theta_j = (1 \times 1)$

### 2.3 Gradient Descent

Hypothesis Function	$h_{\theta}(x) = \theta^T \times x$
Vectorized Hypothesis Function	$h_{\theta}(X) = X \cdot \theta$
Linear Regression Cost Function	$J(\theta) = \frac{1}{2m} \sum (h_{\theta}(X) - y)^2$
Derivative of Linear Regression CF wrt $\theta_j$	$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum (h_{\theta}(X) - y) \cdot X_j$
Change in $\theta_j$	$\begin{aligned} \theta_j &= \theta_j - \alpha \frac{\partial}{\partial \theta_j} \\ &= \theta_j - \alpha \frac{1}{m} \sum (h_{\theta}(X) - y) \cdot X_j \end{aligned}$
Vectorized Change in $\theta$	$\theta = \theta - \alpha \frac{1}{m} X^T (X \cdot \theta - y)$

### 2.4 Normal Equation

$$\theta = (X^T \cdot X)^{-1} \cdot X^T \cdot y$$

### 3 Week 3: Logistic Regression

#### 3.1 Overview

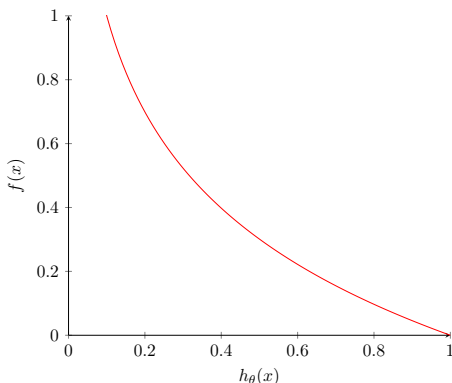
- Use logistic regression for discrete output (classification)
  - $h_{\theta}(x) = (y = 1|x; \theta)$ ; gives probability that the output is 1 given  $x$
  - Sigmoid/Logistic function maps any real number to  $(0, 1)$
  - Logarithm turns sum into product, allowing easier differentiation without altering search space
- For multi-class classification, use one-vs-all
  - Pick class  $i$  that maximizes  $h_{\theta}^i(x)$
- Overfitting is when learned hypothesis fits training data well but fails to generalize; underfitting is when doesn't fit training data
- Address overfitting by reducing number of features, model selection, and regularization
  - Regularization results in simpler hypothesis and less overfitting
  - Extremely large  $\lambda$  will result in underfitting and gradient descent will fail to converge
  - Do not regularize  $\lambda_0$
- Use other prewritten optimization algorithms (conjugate gradient, BFGS, L-BFGS) because they are faster

#### 3.2 Logistic Regression Hypothesis Function

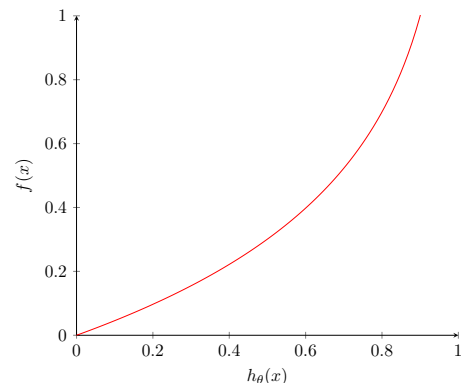
Sigmoid/Logistic Function	$g(z) = \frac{1}{1 + e^{-z}}$
Hypothesis Function	$h_{\theta}(x) = g(\theta^T x)$
	$= \frac{1}{1 + e^{-\theta^T x}}$

#### 3.3 Logistic Regression Cost Function

$$\begin{aligned} \text{Cost}(h_{\theta}(x), y) &= \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases} \\ &= -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x)) \\ J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^i), y^i) \\ J(\theta) &= \frac{-1}{m} \sum_{i=1}^m [y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))] \end{aligned}$$



(a) if  $y=1$



(b) if  $y=0$

### 3.4 Proof of Logistic Regression Cost Function Derivative

$$\begin{aligned}
J(\theta) &= \frac{-1}{m} \sum_{i=1}^m [y^i \log(h_\theta(x^i)) + (1 - y^i) \log(1 - h_\theta(x^i))] \\
\log(h_\theta(x^i)) &= \log\left(\frac{1}{1 + e^{-\theta x^i}}\right) = -\log(1 + e^{-\theta x^i}) \\
\log(1 - h_\theta(x^i)) &= \log\left(1 - \frac{1}{1 + e^{-\theta x^i}}\right) = \log(e^{-\theta x^i}) - \log(1 + e^{-\theta x^i}) = -\theta x^i - \log(1 + e^{-\theta x^i}) \\
J(\theta) &= -\frac{1}{m} \sum_{i=1}^m \left[ -y^i (\log(1 + e^{-\theta x^i})) + (1 - y^i) (-\theta x^i - \log(1 + e^{-\theta x^i})) \right] \\
&= -\frac{1}{m} \sum_{i=1}^m \left[ y^i \theta x^i - \theta x^i - \log(1 + e^{-\theta x^i}) \right] \\
&= -\frac{1}{m} \sum_{i=1}^m \left[ y^i \theta x^i - \log(e^{\theta x^i}) - \log(1 + e^{-\theta x^i}) \right] \\
&= -\frac{1}{m} \sum_{i=1}^m \left[ y^i \theta x^i - \log(1 + e^{\theta x^i}) \right] \\
\frac{\partial}{\partial \theta_j} y^i \theta x^i &= y^i x_j^i \\
\frac{\partial}{\partial \theta_j} \log(1 + e^{\theta x^i}) &= \frac{x_j^i e^{\theta x^i}}{1 + e^{\theta x^i}} \\
&= \frac{x_j^i}{1 + e^{-\theta x^i}} \\
&= x_j^i h_\theta(x^i) \\
\frac{\partial}{\partial \theta_j} J(\theta) &= -\frac{1}{m} \sum_{i=1}^m [y^i x_j^i - x_j^i h_\theta(x^i)] \\
\frac{\partial}{\partial \theta_j} J(\theta) &= \frac{1}{m} \sum_{i=1}^m [h_\theta(x^i) - y^i] x_j^i
\end{aligned}$$

Credit<sup>1</sup>

---

<sup>1</sup><https://math.stackexchange.com/questions/477207/derivative-of-cost-function-for-logistic-regression>

### 3.5 Regularization

Regularizing Term	$\lambda \sum_{j=1}^n \theta_j^2$
Regularized Linear Regression CF	$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^n \theta_j^2$
Regularized Logistic Regression CF	$J(\theta) = \frac{-1}{m} \sum_{i=1}^m [y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$
Regularized GD (Lin/Log Regression)	$\begin{cases} \theta_0 = \theta_j - \alpha \left[ \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) x_0^i \right] \\ \theta_j = \theta_j - \alpha \left[ \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) x_j^i + \frac{\lambda}{m} \theta_j \right] \end{cases} \quad (j=1,2,\dots,n)$
Regularized Normal Equation	$\theta = (X^T X + \lambda \begin{bmatrix} 0 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix}_{n+1,n+1})^{-1} X^T y$

## 4 Week 4: Artificial Neural Networks Representation

### 4.1 Overview

- Neural networks allow for non-linear classification in situations with many features
  - Necessary b/c 100 features at 3rd level polynomials generates 170k features, which quickly becomes intractable
  - “One learning algorithm” hypothesis; you can see with your tongue : brain learns using one algorithm, not thousands of different programs
  - Can have multiple hidden layers
  - Can have multiple outputs (one-vs-all for multi-class classification)
  - If network has  $s_j$  units in layer  $j$  and  $s_{j+1}$  units in layer  $j + 1$ , then  $\Theta^j$  will be of dimension  $s_{j+1} \times s_j + 1$ 
    - The +1 comes from the addition in  $\Theta^{(j)}$  of the bias node,  $x_0$  and  $\Theta_0^{(j)}$
- Forward Propagation is used to predict based on learned parameters
- Bias node gives each node a trainable constant value
  - Allows bias weight to shift the activation curve left/right
  - Other weights affect steepness
- Fun fact: For image recognition, particular order of pixels does not matter for ANN (but does for humans), you just need to keep the convention the same

### 4.2 ANN Notation

- $g(x)$ : sigmoid function
- $\Theta^{(j)}$ : matrix of weights controlling function mapping from layer  $j$  to layer  $j+1$ ; each layer gets own  $\Theta^j$
- $\Theta_{j,0}, \dots, \Theta_{j,n}$  weights corresponding to the inputs  $a_0, \dots, a_n$  going into  $z_j$
- $w_{0,j}, \dots, w_{n,j}$  weights corresponding to the inputs  $a_0, \dots, a_n$  going into  $z_j$
- $z_k^{(j)}$ : encompasses parameters inside of  $g$  function
- $a_i^{(j)}$ : “activation” of unit  $i$  in layer  $j$

### 4.3 ANN Equations

$$\begin{aligned}z_k^{(j)} &= \Theta_{k,0}^{(j-1)}x_0 + \Theta_{k,1}^{(j-1)}x_1 + \dots + \Theta_{k,n}^{(j-1)}x_n \\&= \Theta_{k,0}^{(j-1)}a_0^{(j-1)} + \Theta_{k,1}^{(j-1)}a_1^{(j-1)} + \dots + \Theta_{k,n}^{(j-1)}a_n^{(j-1)} \\z^{(j)} &= \Theta^{(j-1)}a^{(j-1)} \\a^{(j)} &= g(z^{(j)})\end{aligned}$$



#### 4.4 Sample Three Layer System

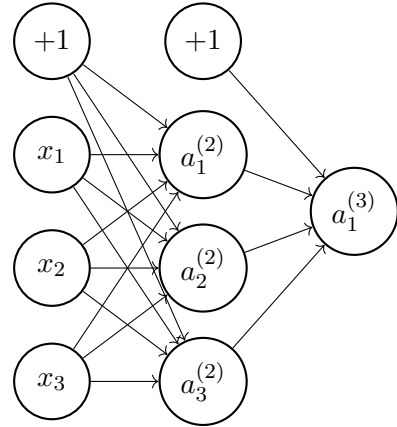
$$\begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \rightarrow \begin{bmatrix} a_1^{(2)} \\ a_2^{(2)} \\ a_3^{(2)} \end{bmatrix} \rightarrow h_{\Theta}(x)$$

$$a_1^{(2)} = g(\Theta_{1,0}^{(1)}x_0 + \Theta_{1,1}^{(1)}x_1 + \Theta_{1,2}^{(1)}x_2 + \Theta_{1,3}^{(1)}x_3)$$

$$a_2^{(2)} = g(\Theta_{2,0}^{(1)}x_0 + \Theta_{2,1}^{(1)}x_1 + \Theta_{2,2}^{(1)}x_2 + \Theta_{2,3}^{(1)}x_3)$$

$$a_3^{(2)} = g(\Theta_{3,0}^{(1)}x_0 + \Theta_{3,1}^{(1)}x_1 + \Theta_{3,2}^{(1)}x_2 + \Theta_{3,3}^{(1)}x_3)$$

$$\begin{aligned} h_{\Theta}(x) &= g(\Theta_{1,0}^{(2)}a_0^{(2)} + \Theta_{1,1}^{(2)}a_1^{(2)} + \Theta_{1,2}^{(2)}a_2^{(2)} + \Theta_{1,3}^{(2)}a_3^{(2)}) \\ &= g(z^3) \\ &= a_1^{(3)} \end{aligned}$$



## 5 Week 5: Artificial Neural Network Learning

### 5.1 Notation

- $L$  : total number of layers in the network
- $s_l$  : number of units (not counting bias unit) in layer  $l$
- $K$  : number of output units/classes
- $h_{\Theta}(x)_k$  : hypothesis that results in the  $k$ th output
- $\delta_k$  : Error signal

### 5.2 Cost Function

$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K \left[ y_k^{(i)} \log((h_{\Theta}(x^{(i)}))_k) + (1 - y_k^{(i)}) \log(1 - (h_{\Theta}(x^{(i)}))_k) \right] + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{j,i}^{(l)})^2$$

Picking training example

Picking output node

Picking layer

Picking node

Picking  $\Theta$

$$\sum_{i=1}^m \sum_{k=1}^K \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}}$$

### 5.3 Backpropagation Algorithm

1. Set  $a(1) := x(t)$
2. Perform forward propagation to compute  $a(l)$  for  $l = 2, 3, \dots, L$
3. Using  $y^{(t)}$ , compute  $\delta^L = a^{(L)} - y^{(t)}$
4. Compute  $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$  using  $\delta^{(l)} = ((\Theta^{(l)})^T \delta^{(l+1)}) \cdot a^{(l)} \cdot (1 - a^{(l)})$
5.  $\Delta_{i,j}^{(l)} := \Delta_{i,j}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$  or with vectorization  $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)} (a^{(l)})^T$
6.  $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = D_{ij}^{(l)}$  where  $\begin{cases} D_{i,j}^{(l)} := \frac{1}{m} \Delta_{i,j}^{(l)} & (j = 0) \\ D_{i,j}^{(l)} := \frac{1}{m} \left( \Delta_{i,j}^{(l)} + \lambda \Theta_{i,j}^{(l)} \right) & (j \neq 0) \end{cases}$

## 5.4 Backpropagation Derivation - Base Case

$$\frac{\partial C}{\partial w_{ij}} = \frac{\partial C}{\partial z_j} \frac{\partial z_j}{\partial w_{ij}} \quad \text{Cost } C \text{ varies wrt input accumulator } z_j, z_j \text{ varies wrt } w_{ij} \text{ (Chain Rule)}$$

$$\frac{\partial C}{\partial z_j} = \frac{\partial}{\partial z_j} (y_j - a_j)^2 \quad \text{Cost defined as } (y_j - a_j)^2$$

$$= \frac{\partial}{\partial z_j} (y_j - g(z_j))^2$$

$$= -2(y_j - g(z_j))g'(z_j)$$

$$g'(z_j) = \frac{d}{dz_j} \frac{1}{1 + e^{-z_j}} \quad \text{Derivative of logistic function}$$

$$= \frac{1}{1 + e^{-z_j}} \frac{e^{-z_j}}{1 + e^{-z_j}}$$

$$= g(z_j)(1 - g(z_j))$$

$$\frac{\partial C}{\partial z_j} = -2(y_j - g(z_j))g(z_j)(1 - g(z_j))$$

$$= -2(y_j - a_j)a_j(1 - a_j)$$

$$\frac{\partial z_j}{\partial w_{ij}} = \frac{\partial}{\partial w_{ij}} \sum_q a_q w_{q,j} \quad \text{Definition of } z_j \text{ as sum of previous node's inputs and their weights}$$

$$= a_i$$

$$\frac{\partial C}{\partial w_{ij}} = -2(y_j - a_j)a_j(1 - a_j)(a_i)$$

$$= -\Delta_j a_i$$

## 5.5 Backpropagation Derivation - Recursive Case

$$\frac{\partial C}{\partial w_{i,j}} = \sum_k \left( \frac{\partial C}{\partial z_k} \frac{\partial z_k}{\partial a_j} \frac{\partial a_j}{\partial z_j} \frac{\partial z_j}{\partial w_{i,j}} \right) \quad C \text{ depends on } z_k, z_k \text{ depends on } a_j, a_j \text{ depends on } z_j, z_j \text{ depends on } w_{i,j}$$

$$\frac{\partial C}{\partial z_k} = -\Delta_k$$

$$\frac{\partial z_k}{\partial a_j} = \frac{\partial}{\partial a_j} \sum_s a_s w_{s,k}$$

$$= w_{j,k}$$

$$\frac{\partial a_j}{\partial z_j} = \frac{\partial}{\partial z_j} g(z_j)$$

$$= g(z_j)(1 - g(z_j))$$

$$= a_j(1 - a_j)$$

$$\frac{\partial z_j}{\partial w_{i,j}} = \frac{\partial}{\partial w_{i,j}} \sum_q a_q w_{q,j}$$

$$= a_i$$

$$\frac{\partial C}{\partial w_{i,j}} = \sum_k (-\Delta_k w_{j,k} a_j (1 - a_j) a_i)$$

$$= \sum_k (-\Delta_k w_{j,k}) a_j (1 - a_j) a_i$$

$$= -\Delta_j a_i$$

## 5.6 Backpropagation Intuition

$\delta_k$  is the error signal from the output

$$\delta_k = (a_k - t_k)g'_k(z_k)$$

so, error wrt output weights is

$$\frac{\partial E}{\partial w_{j,k}} = \delta_k a_j$$

similarly, error wrt internal weights is

$$\frac{\partial E}{\partial w_{i,j}} = \delta_j a_i$$

the error is determined by following layer's error, so it may also be understood as

$$\frac{\partial E}{\partial w_{i,j}} = g'_j(z_j) \sum_k (\delta_k w_{j,k}) a_i$$

## 5.7 Gradient Checking

- Gradient checking ensures that backpropagation is actually working
  - Turn off gradient checking once backpropagation is verified to work
- Approximate the derivative of cost function using slope
  - Pick  $\epsilon = 10^{-4}$
  - Check all  $\Theta_j$

$$\frac{\partial}{\partial \Theta_j} J(\Theta) \approx \frac{J(\Theta_1, \dots, \Theta_j + \epsilon, \dots, \Theta_n) - J(\Theta_1, \dots, \Theta_j - \epsilon, \dots, \Theta_n)}{2\epsilon}$$

## 5.8 Random Initialization

- Need Symmetry Breaking
  - All hidden units would receive the same signal and the same updates
  - Results in finding only one feature, redundantly copied many times over
- Fixed by randomly initializing each  $\Theta_{ij}^{(l)}$  to a value in  $[-\epsilon, \epsilon]$ 
  - A good choice of  $\epsilon$  init is  $\epsilon = \frac{\sqrt{6}}{\sqrt{L_{in} + L_{out}}}$ , where  $L_{in} = s_l$  and  $L_{out} = s_{l+1}$ , the number of units in the layers adjacent to  $\Theta^{(l)}$

<sup>2</sup><https://pdfs.semanticscholar.org/69c2/814c7f8ca16aa836fd32e3c4975f8208e63f.pdf>

<sup>3</sup><https://stats.stackexchange.com/questions/94387/how-to-derive-errors-in-neural-network-with-the-backpropagation-algorithm>

<sup>4</sup><https://theclevermachine.wordpress.com/2014/09/06/derivation-error-backpropagation-gradient-descent-for-neural-networks/>

## 6 Week 6: Advice for Machine Learning

### 6.1 Overview

- Training set: 60%
- Cross validation set: 20%
- Test set: 20%
- Randomize your data and plot learning curves to determine high/low variance

### 6.2 Evaluating Hypothesis

Linear regression

$$J_{test}(\Theta) = \frac{1}{2m_{test}} \sum_{i=1}^{m_{test}} (h_{\Theta}(x_{test}^{(i)}) - y_{test}^{(i)})^2$$

Classification

$$err(h_{\Theta}(x), y) = \begin{cases} 1 & \text{if } h_{\Theta}(x) \geq 0.5 \text{ and } y = 0 \\ 1 & \text{if } h_{\Theta}(x) < 0.5 \text{ and } y = 1 \\ 0 & \text{otherwise} \end{cases}$$

Test Error

$$= \frac{1}{m_{test}} \sum_{i=1}^{m_{test}} err(h_{\Theta}(x_{test}^{(i)}), y_{test}^{(i)})$$

### 6.3 Bias, Variance, Learning Curves

- High variance is when overfitting data
  - Get more training examples, try a smaller set of features, increase lambda
- High bias is when underfitting data
  - Add features, add polynomial features, decrease lambda
- Small neural networks are more prone to underfitting and are cheaper
- Large neural networks are more prone to overfitting and are more expensive

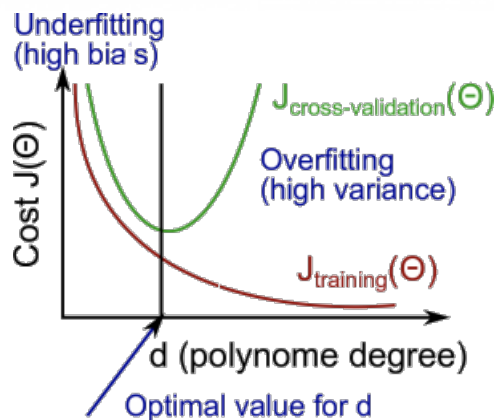
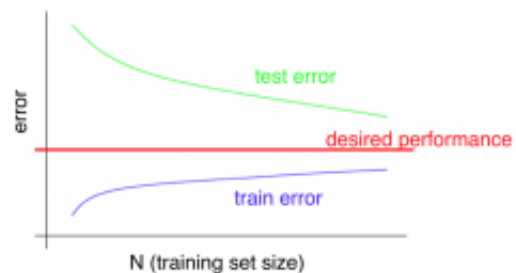
More on Bias vs. Variance

Typical learning curve for high bias (at fixed model complexity):



More on Bias vs. Variance

Typical learning curve for high variance (at fixed model complexity):



## 6.4 Error Analysis

- Difficult to know which feature is actually most helpful
- Start with simple algorithm and test it on cross validation data
- Look at learning curves to decide where to focus
- Manually look at errors in cross validation data to spot trends
- Iterate towards improvement according to a single numerical value of error

## 6.5 Precision and Recall

- Better metric for skewed classes, where there is only a very small chance of  $y = 1$
- Precision =  $\frac{\text{true positives}}{\text{predicted positives}}$
- Recall =  $\frac{\text{true positives}}{\text{actual positives}}$
- Tradeoff between precision and recall by changing threshold
  - Predict 1 if  $h_{\theta}(x) \geq \text{threshold}$
- $F_1$  score =  $2 \frac{PR}{P+R}$ 
  - Maximize  $F_1$  score for picking threshold

## 6.6 Large Data Sets

- Use large data (many training examples) if a human can confidently predict the results from the feature set
- Learning algorithms with many features/hidden units may need more data

## 7 Week 7: Support Vector Machines

### 7.1 Overview

- Change cost function and lambda, but otherwise is basically logistic regression
- Creates larger margins
- Use prewritten SVM's because they have computational tricks to make them efficient enough to work
- Use feature scaling
- Large C: lower bias, higher variance
- Large  $\sigma$ : higher bias, lower variance, varies more smoothly
- If n is small and m is intermediate, use SVM w/ gaussian kernel
- If n is small and m is large, add features and use logistic regression or SVM w/o kernel

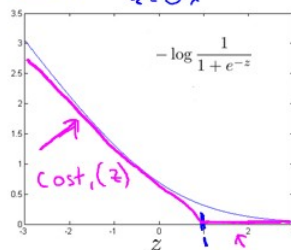
### 7.2 Large Margin Classification

$$\min_{\theta} C \sum_{i=1}^m \left[ y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

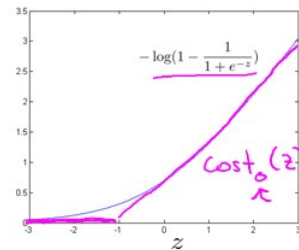
### 7.3 Math behind SVM

- $\theta^T x^{(i)} = p^{(i)} |\theta|$
- Try to maximize  $p^{(i)}$  so that  $|\theta|$  can be minimized

If  $y = 1$  (want  $\theta^T x \gg 0$ ):  
 $z = \theta^T x$



If  $y = 0$  (want  $\theta^T x \ll 0$ ):



### 7.4 Kernels

- Compute features according to proximity to landmarks  $l^{(i)}$ 
  - Pick landmarks based on training data
  - Gaussian Kernel:  $f_1 = \text{similarity}(x, l^{(1)}) = \exp(-\frac{|x-l^{(1)}|^2}{2\sigma^2})$
- Other kernels exist, like linear or polynomial; all must satisfy Mercer's Theorem

## 8 Week 8: Unsupervised Learning

### 8.1 Overview

- Unsupervised learning will find structure in data
- K-Means is a way to cluster
  - Pick number of clusters  $K$  by hand or application
  - K-Means can get stuck in local optima, especially for small numbers of clusters
    - Run algorithm many times, pick cluster with lowest cost

### 8.2 K-Means Notation

- $K$ : number of clusters
- $x^{(1)}$ : training examples  $\in \mathbb{R}$
- $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ : training set
- $\mu_k$ : average of points assigned to cluster  $k$
- $c^{(i)}$ : index of cluster centroid closest to  $x^{(i)}$

### 8.3 K-Means Distortion Function/Optimization Objective

$$\min J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m |x^{(i)} - \mu_{c^{(i)}}|^2$$

### 8.4 K-Means Algorithm

1. Randomly initialized  $K$  cluster centroids according to  $K$  training examples
2.  $c^{(i)} :=$  index of cluster centroid closest to  $x^{(i)}$
3.  $\mu_k :=$  average of points assigned to cluster centroid  $k$

### 8.5 Dimensionality Reduction: Principal Component Analysis

- Useful for data compression, visualizations, and speeding up computation
- PCA projects data onto lower dimensional surface
  - Preprocess data with feature scaling and mean normalization
  - Try to maintain at least 85% variance, preferably 95%+
- Don't use PCA to address overfitting

### 8.6 PCA Notation

- Covariance matrix:  $\Sigma$
- Number of principal components:  $k$
- $[U, S, V] = \text{svd}(\Sigma)$ 
  - Matrix of eigenvectors:  $U$
  - $U_{\text{reduce}} = U(:, 1:k)$
- New inputs:  $z = U_{\text{reduce}}^\top * x$

### 8.7 PCA Equations

$$\Sigma = \frac{1}{m} \sum_{i=1}^n (x^{(i)})(x^{(i)})^\top$$

$$\text{Average squared projection error} = \frac{1}{m} \sum_{i=1}^m |x^{(i)} - x_{\text{approx}}^{(i)}|^2$$



$$\text{Total variance} = \frac{1}{m} \sum_{i=1}^m |x^{(i)}|^2$$

## 8.8 PCA Algorithm

1. Mean normalize and feature scaling
2. Compute covariance matrix  $\Sigma$
3. Compute eigenvectors of  $\Sigma$  using singular value decomposition
4. Compute  $z$

## 8.9 PCA K

$$\min_k \frac{\frac{1}{m} \sum_{i=1}^m |x^{(i)} - x_{approx}|^2}{\frac{1}{m} \sum_{i=1}^m |x^{(i)}|^2} \leq 0.01$$

alternatively, use  $S$  from  $\text{svd}(\Sigma)$

$$\min_k \frac{\sum_{i=1}^k S_{ii}}{\sum_{i=1}^m S_{ii}} \geq 0.99$$

## 9 Week 9: Anomaly Detection

### 9.1 Overview

- Use Anomaly detection over Supervised learning when there are very few positive examples, there are many different types of anomalies, or when future anomalies may not look like current examples
- Gaussian Distribution
  - Positive  $\mu$  shifts bellcurve right
  - Small  $\sigma$  increases height of bellcurve
  - For non-gaussian features, do something like  $x_1 \leftarrow \log(x_1)$
  - Can also create your own features;  $x_3 = \frac{x_1}{x_2}$
  - Choose  $\epsilon$  by maximizing  $F_1$
- Multivariate Gaussian will automatically capture correlations between features
  - Is more computationally expensive bc  $\Sigma^{-1}$
  - Should have  $m > 10n$  to verify invertibility
- For recommender systems, mean normalization can help recommend items for new users

### 9.2 Gaussian (Normal) Distribution Equations

$$p(x) = \prod_{j=1}^n p(x_j; \mu_j; \sigma_j^2) = \prod_{j=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$
$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$
$$\sigma_j^2 = \frac{1}{m} \sum_{i=1}^m (x_j^{(i)} - \mu_j)^2$$

### 9.3 Anomaly Detection Algorithm

1. Choose features  $x_j$  indicative of anomalous examples
2. Fit parameters  $\mu$  and  $\sigma$
3. Compute  $p(x)$ ; anomaly if  $p(x) < \epsilon$

### 9.4 Multivariate Gaussian Distribution Equations

$$p(x; \mu; \Sigma) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(x - \mu)^\top \Sigma^{-1}(x - \mu)\right)$$
$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$
$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)(x^{(i)} - \mu)^\top$$

### 9.5 Recommender Systems Notation

- $n_u$ : number of users
- $n_m$ : number of movies
- $r(i, j) = 1$  if user  $j$  has rated movie  $i$
- $y^{(i,j)}$ : rating given by user  $j$  to movie  $i$
- $\theta^{(j)}$ : parameter vector for user  $j$
- $x^{(i)}$ : feature vector for movie  $i$
- $m^{(j)}$ : number of movies rated by user  $j$

## 9.6 Recommender Systems Algorithm

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} \left( (\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

$$\theta_k^{(j)} = \theta_k^{(j)} - \alpha \sum_{i:r(i,j)=1} \left( (\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right) x_k^{(i)} \quad k=0$$

$$\theta_k^{(j)} = \theta_k^{(j)} - \alpha \left( \sum_{i:r(i,j)=1} \left( (\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right) x_k^{(i)} + \lambda \theta_k^{(j)} \right)$$

## 9.7 Collaborative Filtering Equations

$$J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}) = \frac{1}{2} \sum_{(i,j):r(i,j)=1} \left( (\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{i=1}^{n_m} \sum_{k=1}^n (x_k^{(i)})^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

$$x_k^{(i)} = x_k^{(i)} - \alpha \left( \sum_{i:r(i,j)=1} \left( (\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right) \theta_k^{(j)} + \lambda x_k^{(i)} \right) \quad i=1, \dots, n_m$$

$$\theta_k^{(j)} = \theta_k^{(j)} - \alpha \left( \sum_{i:r(i,j)=1} \left( (\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right) x_k^{(i)} + \lambda \theta_k^{(j)} \right) \quad j=1, \dots, n_u$$

## 9.8 Collaborative Filtering Algorithm

1. Initialize  $x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)}$  to small random values
2. Minimize  $J(x^{(1)}, \dots, x^{(n_m)}, \theta^{(1)}, \dots, \theta^{(n_u)})$
3. Predict rating using  $\theta^\top x$

## 10 Week 10: Large Scale ML

### 10.1 Overview

- Gradient Descent with large datasets
  - Stochastic: 1 example per iteration
  - Batch:  $m$  examples per iteration
  - Mini-Batch:  $b$  examples per iteration;  $b \approx 10$
  - Repeat through set of training examples multiple times if needed
- Online learning
  - Unlimited data because continuous stream, so throw away examples after you use them
  - Can adapt to changing user tastes
- Map-reduce
  - Parallelize by splitting data computation up over multiple computers and recombining later

## 11 Week 11: Photo OCR

### 11.1 Overview

- Problem Description and Pipeline
  - Text Detection, Character Segmentation, Character Classification
- Obtaining lots of data
  - Artificial Data Synthesis
    - Create more data by distorting/modifying current data
  - Get more data by labeling it yourself
  - Crowdsourcing data
    - Amazon Mechanical Turks
- Ceiling Analysis
  - Don't trust gut feeling, give ground truth to pipeline and spend developer time on areas with greatest potential for improvement

### 11.2 Summary

- Supervised Learning
  - Linear regression, logistic regression, neural networks, SVMs
- Unsupervised Learning
  - K-means, PCA, Anomaly detection
- Special Applications / Special topics
  - Recommender systems, large scale machine learning
- Advice on building a machine learning system
  - Bias/variance, regularization; deciding what to work on next; evaluation of learning algorithms, learning curves, error analysis, ceiling analysis