

Coursera-Stanford-ML-Notes

Quentin Truong

20 June 2017 - ? July 2017

Contents

| | | |
|----------|---|-----------|
| 1 | Week 1: Introduction | 2 |
| 1.1 | Overview | 2 |
| 2 | Week 2: Linear Regression with Multiple Variables | 3 |
| 2.1 | Overview | 3 |
| 2.2 | Notation | 3 |
| 2.3 | Gradient Descent | 3 |
| 2.4 | Normal Equation | 3 |
| 3 | Week 3: Logistic Regression | 4 |
| 3.1 | Overview | 4 |
| 3.2 | Logistic Regression Hypothesis Function | 4 |
| 3.3 | Logistic Regression Cost Function | 4 |
| 3.4 | Proof of Logistic Regression Cost Function Derivative | 5 |
| 3.5 | Regularization | 6 |
| 4 | Week 4: Artificial Neural Networks Representation | 7 |
| 4.1 | Overview | 7 |
| 4.2 | ANN Notation | 7 |
| 4.3 | ANN Equations | 7 |
| 4.4 | Sample Three Layer System | 8 |
| 5 | Week 5: Artificial Neural Network Learning | 9 |
| 5.1 | Notation | 9 |
| 5.2 | Cost Function | 9 |
| 5.3 | Backpropagation Algorithm | 9 |
| 5.4 | Backpropagation Derivation - Base Case | 10 |
| 5.5 | Backpropagation Derivation - Recursive Case | 10 |
| 5.6 | Backpropagation Intuition | 11 |
| 5.7 | Gradient Checking | 11 |
| 5.8 | Random Initialization | 11 |
| 6 | Week 6: Advice for Machine Learning | 12 |
| 6.1 | Overview | 12 |
| 6.2 | Evaluating Hypothesis | 12 |
| 6.3 | Bias, Variance, Learning Curves | 12 |
| 6.4 | Error Analysis | 13 |
| 6.5 | Precision and Recall | 13 |
| 6.6 | Large Data Sets | 13 |

| | | |
|----------|--|-----------|
| 7 | Week 7: Support Vector Machines | 14 |
| 7.1 | Overview | 14 |
| 7.2 | Large Margin Classification | 14 |
| 7.3 | Math behind SVM | 14 |
| 7.4 | Kernels | 14 |
| 8 | Week 8: Unsupervised Learning | 15 |
| 8.1 | Overview | 15 |
| 8.2 | K-Means Notation | 15 |
| 8.3 | K-Means Distortion Function/Optimization Objective | 15 |
| 8.4 | K-Means Algorithm | 15 |
| 8.5 | Dimensionality Reduction: Principal Component Analysis | 15 |
| 8.6 | PCA Notation | 15 |
| 8.7 | PCA: Equations | 15 |
| 8.8 | PCA Algorithm | 16 |
| 8.9 | PCA K | 16 |
| 9 | Week 9: Anomaly Detection | 17 |
| 9.1 | Overview | 17 |
| 9.2 | Gaussian (Normal) Distribution | 17 |
| 9.3 | Anomaly Detection Algorithm | 17 |
| 9.4 | Multivariate Gaussian Distribution | 17 |
| 9.5 | Recommender Systems Notation | 17 |
| 9.6 | Recommender Systems Algorithm | 18 |

1 Week 1: Introduction

1.1 Overview

- Machine Learning: “A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .”
- Supervised Learning: know what our correct output looks like
 - Regression: want continuous output
 - Classification: want discrete output
- Unsupervised Learning: little or no idea what our results should look like
 - Clustering: find groups according to similarity in various variables
 - Nonclustering: find structure in chaos

2 Week 2: Linear Regression with Multiple Variables

2.1 Overview

- Use linear regression for continuous output
- Choose gradient descent if many features (million+) because the inverse matrix required for the normal equation can become expensive to compute
- Normal equation will directly compute theta
- Normalize features if using gradient descent

2.2 Notation

$m = \text{number of samples}$

$n = \text{number of feature}$

$x = (n \times 1)$

$X = (m \times n)$

$X_j = (m \times 1)$

$\theta = (n \times 1)$

$\theta_j = (1 \times 1)$

2.3 Gradient Descent

| | |
|---|---|
| Hypothesis Function | $h_{\theta}(x) = \theta^T \times x$ |
| Vectorized Hypothesis Function | $h_{\theta}(X) = X \cdot \theta$ |
| Linear Regression Cost Function | $J(\theta) = \frac{1}{2m} \sum (h_{\theta}(X) - y)^2$ |
| Derivative of Linear Regression CF wrt θ_j | $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum (h_{\theta}(X) - y) \cdot X_j$ |
| Change in θ_j | $\begin{aligned} \theta_j &= \theta_j - \alpha \frac{\partial}{\partial \theta_j} \\ &= \theta_j - \alpha \frac{1}{m} \sum (h_{\theta}(X) - y) \cdot X_j \end{aligned}$ |
| Vectorized Change in θ | $\theta = \theta - \alpha \frac{1}{m} X^T (X \cdot \theta - y)$ |

2.4 Normal Equation

$$\theta = (X^T \cdot X)^{-1} \cdot X^T \cdot y$$

3 Week 3: Logistic Regression

3.1 Overview

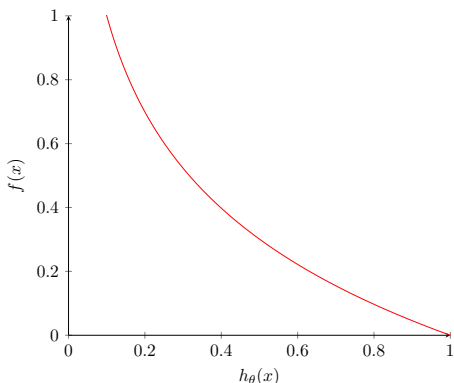
- Use logistic regression for discrete output (classification)
 - $h_{\theta}(x) = (y = 1|x; \theta)$; gives probability that the output is 1 given x
 - Sigmoid/Logistic function maps any real number to $(0, 1)$
 - Logarithm turns sum into product, allowing easier differentiation without altering search space
- For multi-class classification, use one-vs-all
 - Pick class i that maximizes $h_{\theta}^i(x)$
- Overfitting is when learned hypothesis fits training data well but fails to generalize; underfitting is when doesn't fit training data
- Address overfitting by reducing number of features, model selection, and regularization
 - Regularization results in simpler hypothesis and less overfitting
 - Extremely large λ will result in underfitting and gradient descent will fail to converge
 - Do not regularize λ_0
- Use other prewritten optimization algorithms (conjugate gradient, BFGS, L-BFGS) because they are faster

3.2 Logistic Regression Hypothesis Function

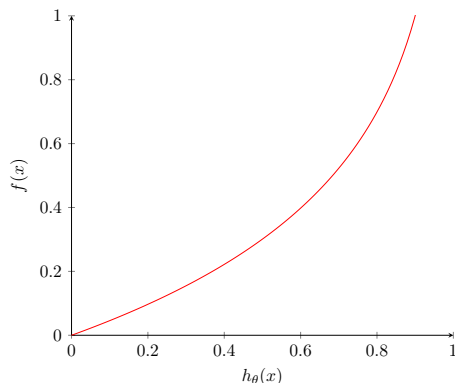
| | |
|---------------------------|-----------------------------------|
| Sigmoid/Logistic Function | $g(z) = \frac{1}{1 + e^{-z}}$ |
| Hypothesis Function | $h_{\theta}(x) = g(\theta^T x)$ |
| | $= \frac{1}{1 + e^{-\theta^T x}}$ |

3.3 Logistic Regression Cost Function

$$\begin{aligned} \text{Cost}(h_{\theta}(x), y) &= \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases} \\ &= -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x)) \\ J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^i), y^i) \\ J(\theta) &= \frac{-1}{m} \sum_{i=1}^m [y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))] \end{aligned}$$



(a) if $y=1$



(b) if $y=0$

3.4 Proof of Logistic Regression Cost Function Derivative

$$\begin{aligned}
J(\theta) &= \frac{-1}{m} \sum_{i=1}^m [y^i \log(h_\theta(x^i)) + (1 - y^i) \log(1 - h_\theta(x^i))] \\
\log(h_\theta(x^i)) &= \log\left(\frac{1}{1 + e^{-\theta x^i}}\right) = -\log(1 + e^{-\theta x^i}) \\
\log(1 - h_\theta(x^i)) &= \log\left(1 - \frac{1}{1 + e^{-\theta x^i}}\right) = \log(e^{-\theta x^i}) - \log(1 + e^{-\theta x^i}) = -\theta x^i - \log(1 + e^{-\theta x^i}) \\
J(\theta) &= -\frac{1}{m} \sum_{i=1}^m \left[-y^i (\log(1 + e^{-\theta x^i})) + (1 - y^i) (-\theta x^i - \log(1 + e^{-\theta x^i})) \right] \\
&= -\frac{1}{m} \sum_{i=1}^m \left[y^i \theta x^i - \theta x^i - \log(1 + e^{-\theta x^i}) \right] \\
&= -\frac{1}{m} \sum_{i=1}^m \left[y^i \theta x^i - \log(e^{\theta x^i}) - \log(1 + e^{-\theta x^i}) \right] \\
&= -\frac{1}{m} \sum_{i=1}^m \left[y^i \theta x^i - \log(1 + e^{\theta x^i}) \right] \\
\frac{\partial}{\partial \theta_j} y^i \theta x^i &= y^i x_j^i \\
\frac{\partial}{\partial \theta_j} \log(1 + e^{\theta x^i}) &= \frac{x_j^i e^{\theta x^i}}{1 + e^{\theta x^i}} \\
&= \frac{x_j^i}{1 + e^{-\theta x^i}} \\
&= x_j^i h_\theta(x^i) \\
\frac{\partial}{\partial \theta_j} J(\theta) &= -\frac{1}{m} \sum_{i=1}^m [y^i x_j^i - x_j^i h_\theta(x^i)] \\
\frac{\partial}{\partial \theta_j} J(\theta) &= \frac{1}{m} \sum_{i=1}^m [h_\theta(x^i) - y^i] x_j^i
\end{aligned}$$

Credit¹

¹<https://math.stackexchange.com/questions/477207/derivative-of-cost-function-for-logistic-regression>

3.5 Regularization

| | |
|-------------------------------------|---|
| Regularizing Term | $\lambda \sum_{j=1}^n \theta_j^2$ |
| Regularized Linear Regression CF | $J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2 + \lambda \sum_{j=1}^n \theta_j^2$ |
| Regularized Logistic Regression CF | $J(\theta) = \frac{-1}{m} \sum_{i=1}^m [y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$ |
| Regularized GD (Lin/Log Regression) | $\begin{cases} \theta_0 = \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) x_0^i \right] \\ \theta_j = \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) x_j^i + \frac{\lambda}{m} \theta_j \right] \quad (j=1,2,\dots,n) \end{cases}$ |
| Regularized Normal Equation | $\theta = (X^T X + \lambda \begin{bmatrix} 0 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix}_{n+1,n+1})^{-1} X^T y$ |

4 Week 4: Artificial Neural Networks Representation

4.1 Overview

- Neural networks allow for non-linear classification in situations with many features
 - Necessary b/c 100 features at 3rd level polynomials generates 170k features, which quickly becomes intractable
 - “One learning algorithm” hypothesis; you can see with your tongue : brain learns using one algorithm, not thousands of different programs
 - Can have multiple hidden layers
 - Can have multiple outputs (one-vs-all for multi-class classification)
 - If network has s_j units in layer j and s_{j+1} units in layer $j + 1$, then Θ^j will be of dimension $s_{j+1} \times s_j + 1$
 - The +1 comes from the addition in $\Theta^{(j)}$ of the bias node, x_0 and $\Theta_0^{(j)}$
- Forward Propagation is used to predict based on learned parameters
- Bias node gives each node a trainable constant value
 - Allows bias weight to shift the activation curve left/right
 - Other weights affect steepness
- Fun fact: For image recognition, particular order of pixels does not matter for ANN (but does for humans), you just need to keep the convention the same

4.2 ANN Notation

- $g(x)$: sigmoid function
- $\Theta^{(j)}$: matrix of weights controlling function mapping from layer j to layer $j+1$; each layer gets own Θ^j
- $\Theta_{j,0}, \dots, \Theta_{j,n}$ weights corresponding to the inputs a_0, \dots, a_n going into z_j
- $w_{0,j}, \dots, w_{n,j}$ weights corresponding to the inputs a_0, \dots, a_n going into z_j
- $z_k^{(j)}$: encompasses parameters inside of g function
- $a_i^{(j)}$: “activation” of unit i in layer j

4.3 ANN Equations

$$\begin{aligned}z_k^{(j)} &= \Theta_{k,0}^{(j-1)} x_0 + \Theta_{k,1}^{(j-1)} x_1 + \dots + \Theta_{k,n}^{(j-1)} x_n \\&= \Theta_{k,0}^{(j-1)} a_0^{(j-1)} + \Theta_{k,1}^{(j-1)} a_1^{(j-1)} + \dots + \Theta_{k,n}^{(j-1)} a_n^{(j-1)} \\z^{(j)} &= \Theta^{(j-1)} a^{(j-1)} \\a^{(j)} &= g(z^{(j)})\end{aligned}$$

4.4 Sample Three Layer System

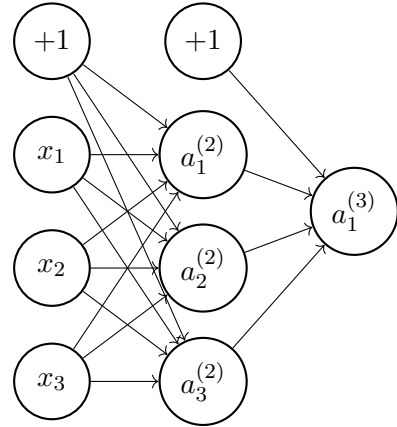
$$\begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \rightarrow \begin{bmatrix} a_1^{(2)} \\ a_2^{(2)} \\ a_3^{(2)} \end{bmatrix} \rightarrow h_{\Theta}(x)$$

$$a_1^{(2)} = g(\Theta_{1,0}^{(1)}x_0 + \Theta_{1,1}^{(1)}x_1 + \Theta_{1,2}^{(1)}x_2 + \Theta_{1,3}^{(1)}x_3)$$

$$a_2^{(2)} = g(\Theta_{2,0}^{(1)}x_0 + \Theta_{2,1}^{(1)}x_1 + \Theta_{2,2}^{(1)}x_2 + \Theta_{2,3}^{(1)}x_3)$$

$$a_3^{(2)} = g(\Theta_{3,0}^{(1)}x_0 + \Theta_{3,1}^{(1)}x_1 + \Theta_{3,2}^{(1)}x_2 + \Theta_{3,3}^{(1)}x_3)$$

$$\begin{aligned} h_{\Theta}(x) &= g(\Theta_{1,0}^{(2)}a_0^{(2)} + \Theta_{1,1}^{(2)}a_1^{(2)} + \Theta_{1,2}^{(2)}a_2^{(2)} + \Theta_{1,3}^{(2)}a_3^{(2)}) \\ &= g(z^3) \\ &= a_1^{(3)} \end{aligned}$$



5 Week 5: Artificial Neural Network Learning

5.1 Notation

- L : total number of layers in the network
- s_l : number of units (not counting bias unit) in layer l
- K : number of output units/classes
- $h_{\Theta}(x)_k$: hypothesis that results in the k th output
- δ_k : Error signal

5.2 Cost Function

$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K \left[y_k^{(i)} \log((h_{\Theta}(x^{(i)}))_k) + (1 - y_k^{(i)}) \log(1 - (h_{\Theta}(x^{(i)}))_k) \right] + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{j,i}^{(l)})^2$$

Picking training example

Picking output node

Picking layer

Picking node

Picking Θ

$$\sum_{i=1}^m \sum_{k=1}^K \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}}$$

5.3 Backpropagation Algorithm

1. Set $a(1) := x(t)$
2. Perform forward propagation to compute $a(l)$ for $l = 2, 3, \dots, L$
3. Using $y^{(t)}$, compute $\delta^L = a^{(L)} - y^{(t)}$
4. Compute $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$ using $\delta^{(l)} = ((\Theta^{(l)})^T \delta^{(l+1)}) \cdot * a^{(l)} \cdot * (1 - a^{(l)})$
5. $\Delta_{i,j}^{(l)} := \Delta_{i,j}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$ or with vectorization $\Delta^{(l)} := \Delta^{(l)} + \delta^{(l+1)}(a^{(l)})^T$
6. $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = D_{ij}^{(l)}$ where $\begin{cases} D_{i,j}^{(l)} := \frac{1}{m} \Delta_{i,j}^{(l)} & (j = 0) \\ D_{i,j}^{(l)} := \frac{1}{m} \left(\Delta_{i,j}^{(l)} + \lambda \Theta_{i,j}^{(l)} \right) & (j \neq 0) \end{cases}$

5.4 Backpropagation Derivation - Base Case

$$\begin{aligned}
\frac{\partial C}{\partial w_{ij}} &= \frac{\partial C}{\partial z_j} \frac{\partial z_j}{\partial w_{ij}} && \text{Cost } C \text{ varies wrt input accumulator } z_j, z_j \text{ varies wrt } w_{ij} \text{ (Chain Rule)} \\
\frac{\partial C}{\partial z_j} &= \frac{\partial}{\partial z_j} (y_j - a_j)^2 && \text{Cost defined as } (y_j - a_j)^2 \\
&= \frac{\partial}{\partial z_j} (y_j - g(z_j))^2 \\
&= -2(y_j - g(z_j))g'(z_j) \\
g'(z_j) &= \frac{d}{dz_j} \frac{1}{1 + e^{-z_j}} && \text{Derivative of logistic function} \\
&= \frac{1}{1 + e^{-z_j}} \frac{e^{-z_j}}{1 + e^{-z_j}} \\
&= g(z_j)(1 - g(z_j)) \\
\frac{\partial C}{\partial z_j} &= -2(y_j - g(z_j))g(z_j)(1 - g(z_j)) \\
&= -2(y_j - a_j)a_j(1 - a_j) \\
\frac{\partial z_j}{\partial w_{ij}} &= \frac{\partial}{\partial w_{ij}} \sum_q a_q w_{q,j} && \text{Definition of } z_j \text{ as sum of previous node's inputs and their weights} \\
&= a_i \\
\frac{\partial C}{\partial w_{ij}} &= -2(y_j - a_j)a_j(1 - a_j)(a_i) \\
&= -\Delta_j a_i
\end{aligned}$$

5.5 Backpropagation Derivation - Recursive Case

$$\begin{aligned}
\frac{\partial C}{\partial w_{i,j}} &= \sum_k \left(\frac{\partial C}{\partial z_k} \frac{\partial z_k}{\partial a_j} \frac{\partial a_j}{\partial z_j} \frac{\partial z_j}{\partial w_{i,j}} \right) && C \text{ depends on } z_k, z_k \text{ depends on } a_j, a_j \text{ depends on } z_j, z_j \text{ depends on } w_{i,j} \\
\frac{\partial C}{\partial z_k} &= -\Delta_k \\
\frac{\partial z_k}{\partial a_j} &= \frac{\partial}{\partial a_j} \sum_s a_s w_{s,k} \\
&= w_{j,k} \\
\frac{\partial a_j}{\partial z_j} &= \frac{\partial}{\partial z_j} g(z_j) \\
&= g(z_j)(1 - g(z_j)) \\
&= a_j(1 - a_j) \\
\frac{\partial z_j}{\partial w_{i,j}} &= \frac{\partial}{\partial w_{i,j}} \sum_q a_q w_{q,j} \\
&= a_i \\
\frac{\partial C}{\partial w_{i,j}} &= \sum_k (-\Delta_k w_{j,k} a_j(1 - a_j) a_i) \\
&= \sum_k (-\Delta_k w_{j,k}) a_j(1 - a_j) a_i \\
&= -\Delta_j a_i
\end{aligned}$$

5.6 Backpropagation Intuition

δ_k is the error signal from the output

$$\delta_k = (a_k - t_k)g'_k(z_k)$$

so, error wrt output weights is

$$\frac{\partial E}{\partial w_{j,k}} = \delta_k a_j$$

similarly, error wrt internal weights is

$$\frac{\partial E}{\partial w_{i,j}} = \delta_j a_i$$

the error is determined by following layer's error, so it may also be understood as

$$\frac{\partial E}{\partial w_{i,j}} = g'_j(z_j) \sum_k (\delta_k w_{j,k}) a_i$$

5.7 Gradient Checking

- Gradient checking ensures that backpropagation is actually working
 - Turn off gradient checking once backpropagation is verified to work
- Approximate the derivative of cost function using slope
 - Pick $\epsilon = 10^{-4}$
 - Check all Θ_j

$$\frac{\partial}{\partial \Theta_j} J(\Theta) \approx \frac{J(\Theta_1, \dots, \Theta_j + \epsilon, \dots, \Theta_n) - J(\Theta_1, \dots, \Theta_j - \epsilon, \dots, \Theta_n)}{2\epsilon}$$

5.8 Random Initialization

- Need Symmetry Breaking
 - All hidden units would receive the same signal and the same updates
 - Results in finding only one feature, redundantly copied many times over
- Fixed by randomly initializing each $\Theta_{ij}^{(l)}$ to a value in $[-\epsilon, \epsilon]$
 - A good choice of ϵ init is $\epsilon = \frac{\sqrt{6}}{\sqrt{L_{in} + L_{out}}}$, where $L_{in} = s_l$ and $L_{out} = s_{l+1}$, the number of units in the layers adjacent to $\Theta^{(l)}$

²<https://pdfs.semanticscholar.org/69c2/814c7f8ca16aa836fd32e3c4975f8208e63f.pdf>

³<https://stats.stackexchange.com/questions/94387/how-to-derive-errors-in-neural-network-with-the-backpropagation-algorithm>

⁴<https://theclevermachine.wordpress.com/2014/09/06/derivation-error-backpropagation-gradient-descent-for-neural-networks/>

6 Week 6: Advice for Machine Learning

6.1 Overview

- Training set: 60%
- Cross validation set: 20%
- Test set: 20%
- Randomize your data and plot learning curves to determine high/low variance

6.2 Evaluating Hypothesis

| | |
|-------------------|---|
| Linear regression | $J_{test}(\Theta) = \frac{1}{2m_{test}} \sum_{i=1}^{m_{test}} (h_{\Theta}(x_{test}^{(i)}) - y_{test}^{(i)})^2$ |
| Classification | $err(h_{\Theta}(x), y) = \begin{cases} 1 & \text{if } h_{\Theta}(x) \geq 0.5 \text{ and } y = 0 \\ 1 & \text{if } h_{\Theta}(x) < 0.5 \text{ and } y = 1 \\ 0 & \text{otherwise} \end{cases}$ |
| Test Error | $= \frac{1}{m_{test}} \sum_{i=1}^{m_{test}} err(h_{\Theta}(x_{test}^{(i)}), y_{test}^{(i)})$ |

6.3 Bias, Variance, Learning Curves

- High variance is when overfitting data
 - Get more training examples, try a smaller set of features, increase lambda
- High bias is when underfitting data
 - Add features, add polynomial features, decrease lambda
- Small neural networks are more prone to underfitting and are cheaper
- Large neural networks are more prone to overfitting and are more expensive


More on Bias vs. Variance

Typical learning curve for high bias (at fixed model complexity):



More on Bias vs. Variance

Typical learning curve for high variance (at fixed model complexity):



6.4 Error Analysis

- Difficult to know which feature is actually most helpful
- Start with simple algorithm and test it on cross validation data
- Look at learning curves to decide where to focus
- Manually look at errors in cross validation data to spot trends
- Iterate towards improvement according to a single numerical value of error

6.5 Precision and Recall

- Better metric for skewed classes, where there is only a very small chance of $y = 1$
- Precision = $\frac{\text{true positives}}{\text{predicted positives}}$
- Recall = $\frac{\text{true positives}}{\text{actual positives}}$
- Tradeoff between precision and recall by changing threshold
 - Predict 1 if $h_{\theta}(x) \geq \text{threshold}$
- F_1 score = $2 \frac{PR}{P+R}$
 - Maximize F_1 score for picking threshold

6.6 Large Data Sets

- Use large data (many training examples) if a human can confidently predict the results from the feature set
- Learning algorithms with many features/hidden units may need more data

7 Week 7: Support Vector Machines

7.1 Overview

- Change cost function and lambda, but otherwise is basically logistic regression
- Creates larger margins
- Use prewritten SVM's because they have computational tricks to make them efficient enough to work
- Use feature scaling
- Large C: lower bias, higher variance
- Large σ : higher bias, lower variance, varies more smoothly
- If n is small and m is intermediate, use SVM w/ gaussian kernel
- If n is small and m is large, add features and use logistic regression or SVM w/o kernel

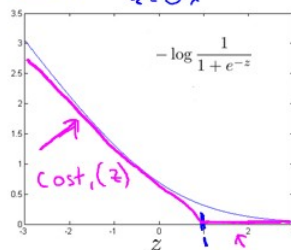
7.2 Large Margin Classification

$$\min_{\theta} C \sum_{i=1}^m \left[y^{(i)} \text{cost}_1(\theta^T x^{(i)}) + (1 - y^{(i)}) \text{cost}_0(\theta^T x^{(i)}) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2$$

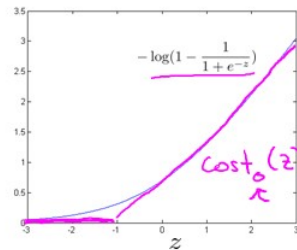
7.3 Math behind SVM

- $\theta^T x^{(i)} = p^{(i)} |\theta|$
- Try to maximize $p^{(i)}$ so that $|\theta|$ can be minimized

If $y = 1$ (want $\theta^T x \gg 0$):
 $z = \theta^T x$



If $y = 0$ (want $\theta^T x \ll 0$):



7.4 Kernels

- Compute features according to proximity to landmarks $l^{(i)}$
 - Pick landmarks based on training data
 - Gaussian Kernel: $f_1 = \text{similarity}(x, l^{(1)}) = \exp(-\frac{|x-l^{(1)}|^2}{2\sigma^2})$
- Other kernels exist, like linear or polynomial; all must satisfy Mercer's Theorem

8 Week 8: Unsupervised Learning

8.1 Overview

- Unsupervised learning will find structure in data
- K-Means is a way to cluster
 - Pick number of clusters K by hand or application
 - K-Means can get stuck in local optima, especially for small numbers of clusters
 - Run algorithm many times, pick cluster with lowest cost

8.2 K-Means Notation

- K : number of clusters
- $x^{(1)}$: training examples $\in \mathbb{R}$
- $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$: training set
- μ_k : average of points assigned to cluster k
- $c^{(i)}$: index of cluster centroid closest to $x^{(i)}$

8.3 K-Means Distortion Function/Optimization Objective

$$\min J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m |x^{(i)} - \mu_{c^{(i)}}|^2$$

8.4 K-Means Algorithm

1. Randomly initialized K cluster centroids according to K training examples
2. $c^{(i)} :=$ index of cluster centroid closest to $x^{(i)}$
3. $\mu_k :=$ average of points assigned to cluster centroid k

8.5 Dimensionality Reduction: Principal Component Analysis

- Useful for data compression, visualizations, and speeding up computation
- PCA projects data onto lower dimensional surface
 - Preprocess data with feature scaling and mean normalization
 - Try to maintain at least 85% variance, preferably 95%+
- Don't use PCA to address overfitting

8.6 PCA Notation

- Covariance matrix: Σ
- Number of principal components: k
- $[U, S, V] = \text{svd}(\Sigma)$
 - Matrix of eigenvectors: U
 - $U_{\text{reduce}} = U(:, 1:k)$
- New inputs: $z = U_{\text{reduce}}^\top * x$

8.7 PCA: Equations

$$\Sigma = \frac{1}{m} \sum_{i=1}^n (x^{(i)})(x^{(i)})^\top$$

$$\text{Average squared projection error} = \frac{1}{m} \sum_{i=1}^m |x^{(i)} - x_{\text{approx}}^{(i)}|^2$$

$$\text{Total variance} = \frac{1}{m} \sum_{i=1}^m |x^{(i)}|^2$$

8.8 PCA Algorithm

1. Mean normalize and feature scaling
2. Compute covariance matrix Σ
3. Compute eigenvectors of Σ using singular value decomposition
4. Compute z

8.9 PCA K

$$\min_k \frac{\frac{1}{m} \sum_{i=1}^m |x^{(i)} - x_{approx}|^2}{\frac{1}{m} \sum_{i=1}^m |x^{(i)}|^2} \leq 0.01$$

alternatively, use S from $\text{svd}(\Sigma)$

$$\min_k \frac{\sum_{i=1}^k S_{ii}}{\sum_{i=1}^m S_{ii}} \geq 0.99$$

9 Week 9: Anomaly Detection

9.1 Overview

- Use Anomaly detection over Supervised learning when there are a very few positive examples, there are many different types of anomalies, or when future anomalies may not look like current examples
- Gaussian Distribution
 - Positive μ shifts bellcurve right
 - Small σ increases height of bellcurve
 - For non-gaussian features, do something like $x_1 \leftarrow \log(x_1)$
 - Can also create your own features; $x_3 = \frac{x_1}{x_2}$
- Multivariate Gaussian will automatically capture correlations between features
 - Is more computationally expensive bc Σ^{-1}
 - Should have $m > 10n$ to verify invertibility

9.2 Gaussian (Normal) Distribution

$$p(x) = \prod_{j=1}^n p(x_j; \mu_j; \sigma_j^2) = \prod_{j=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$
$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$
$$\sigma_j^2 = \frac{1}{m} \sum_{i=1}^m (x_j^{(i)} - \mu_j)^2$$

9.3 Anomaly Detection Algorithm

1. Choose features x_j indicative of anomalous examples
2. Fit parameters μ and σ
3. Compute $p(x)$; anomaly if $p(x) < \epsilon$

9.4 Multivariate Gaussian Distribution

$$p(x; \mu; \Sigma) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(x - \mu)^\top \Sigma^{-1}(x - \mu)\right)$$
$$\mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$
$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)(x^{(i)} - \mu)^\top$$

9.5 Recommender Systems Notation

- n_u : number of users
- n_m : number of movies
- $r(i, j) = 1$ if user j has rated movie i
- $y^{(i,j)}$: rating given by user j to movie i
- $\theta^{(j)}$: parameter vector for user j
- $x^{(i)}$: feature vector for movie i
- $m^{(j)}$: number of movies rated by user j

9.6 Recommender Systems Algorithm

$$\min_{\theta^{(1)}, \dots, \theta^{(n_u)}} \frac{1}{2} \sum_{j=1}^{n_u} \sum_{i:r(i,j)=1} \left((\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right)^2 + \frac{\lambda}{2} \sum_{j=1}^{n_u} \sum_{k=1}^n (\theta_k^{(j)})^2$$

$$\theta_k^{(j)} = \theta_k^{(j)} - \alpha \sum_{i:r(i,j)=1} \left((\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right) x_k^{(i)} \quad k=0$$

$$\theta_k^{(j)} = \theta_k^{(j)} - \alpha \left(\sum_{i:r(i,j)=1} \left((\theta^{(j)})^\top x^{(i)} - y^{(i,j)} \right) x_k^{(i)} + \lambda \theta_k^{(j)} \right)$$